

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
```

```
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE * f;
```

```
    for(i=0; i<MAXPAROLA; i++)
        freq[i]=0;
```

```
    if(argc != 2)
```

```
    {
        fprintf(stderr, "ERRORE: serve un parametro con il nome del file\n");
        exit(1);
    }
```

```
    f = fopen(argv[1], "rt");
    if(f==NULL)
```

```
    {
        fprintf(stderr, "ERRORE: impossibile aprire il file %s\n", argv[1]);
        exit(1);
    }
```

```
    while( fgets( riga, MAXRIGA, f ) != NULL )
```

# Threads

## Threads

Stefano Quer and Pietro Laface

Dipartimento di Automatica e Informatica

Politecnico di Torino

## Processes: Characteristics

- ❖ A process may execute other processes through
  - Cloning (UNIX, **fork**)
  - Replacing the current image with another image (UNIX, **exec**)
  - explicit call (Windows, **CreateProcess**)
- ❖ A process has
  - Its own address space
  - A single execution thread (a single program counter)

## Processes: Limits

- ❖ Synchronization and data transfer
  - No cost or minimal for (almost) independent processes
  - High cost for cooperating processes
- ❖ Cloning involves
  - A significant increase in the memory used
  - Creation time overhead
- ❖ Management of multiple processes requires
  - Scheduling
  - Expensive context switching operations

Standard process = **heavyweight process**  
A task with a single thread of execution

## From processes to threads

- ❖ There are several cases where it would be useful to have
  - Lower creation and management costs
  - A single address space
  - Multiple execution threads within that address space
- ❖ Example
  - WEB applications
    - A server must respond quickly to many access requests
    - The requests are submitted at the same time, and require similar processing of the data, etc.

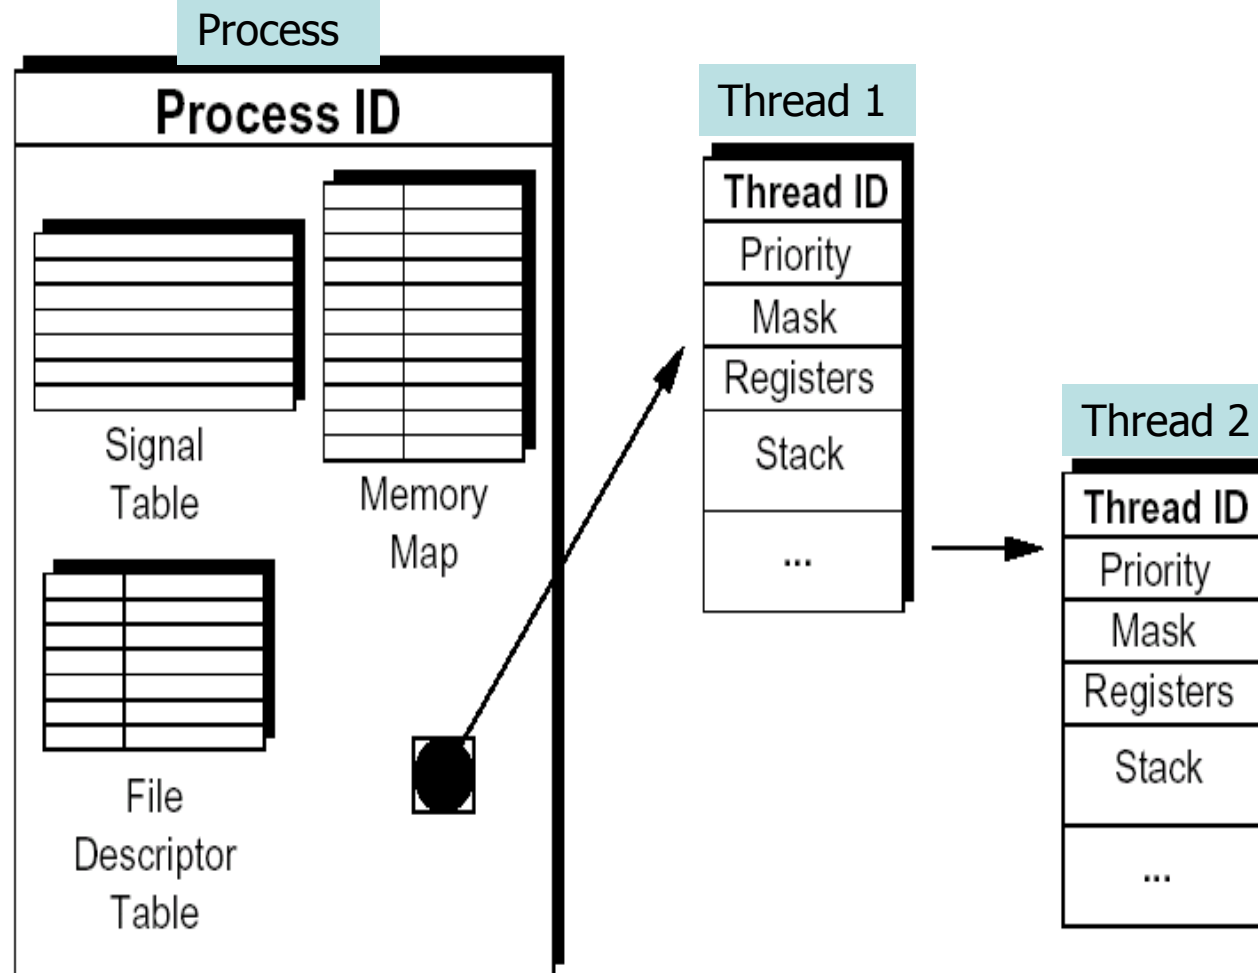
## From processes to threads

- ❖ The 1003.1c POSIX standard introduces the concept of **threads**
- ❖ The thread model allows a program to control multiple different flows of operations that overlap in time.
- ❖ Each flow of operations is referred to as a **thread**
- ❖ Creation and control over these flows is achieved by making calls to the POSIX Threads API.
- ❖ A thread can share its address space with other threads

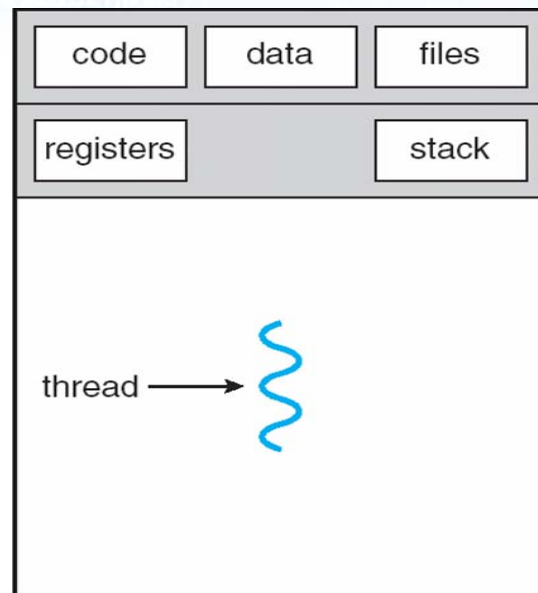
## From processes to threads

- ❖ The **process** is the owner of the resources that are used by all its threads
- ❖ The **thread** is the basic unit of CPU utilization (and scheduling)
- ❖ Thread is also called a **lightweight process**

# From processes to threads

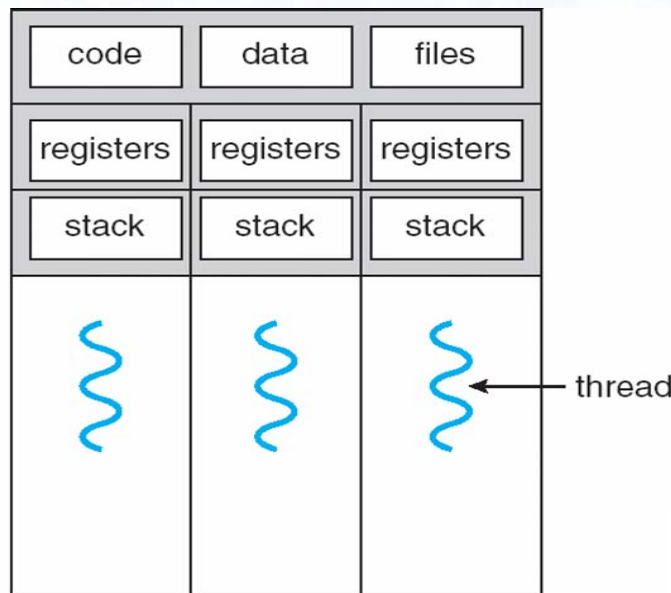


# From processes to threads



single-threaded process

A process with a  
single thread



multithreaded process

A process with three threads  
**Sharing requires protection !**



## Threads: Pros

### ❖ The use of threads allows

#### ➤ Shorter response time

- Creating a thread it is 10-100 times faster than creating a process
- Example
  - To create 50000 jobs (`fork`) takes 10 seconds (real time)
  - To create 50000 thread (`pthread_create`) takes 0.8 seconds (real time)

#### ➤ Shared resources

- Processes can share data only with special techniques
- Threads share data automatically

## Threads: Pros

### ➤ Lower costs for resource management

- Allocate memory to a process is expensive
- Threads use the same section of code and/or data to serve more clients

### ➤ Increased scalability

- The advantages of multi-threaded programming increase in multi-processor systems
- In multi-core systems (different calculation units per processor) threads allow easily implementing concurrent programming paradigms based on
  - Task separation (pipelining)
  - Data partitioning (same task on data blocks)

## Threads: Cons

- ❖ Since threads of the same process run in the same address space they must be synchronized to properly access shared data

## Concurrency with threads

- ❖ Optimize the following code segment that performs the scalar product of two huge dimension vectors, which have been independently created by two processes or by two threads

```
for (i=0; i<n; i++) {  
    v[i] = v1[i] * v2[i];  
}
```

- ❖ Sharing data would be ineffective for processes, but is inexpensive for threads

## Concurrency with threads

```
mult (a, b) {  
    for (i=a; i<b; i++)  
        v[i] = v1[i] * v2[i]  
}  
...  
CreateThread (mult, 0, n/2);  
CreateThread (mult, n/2, n);
```

Data partition with a divide-and-conquer strategy

A thread perform its task on its partition of the data

Care has to be taken to avoid

- the use of non-reentrant procedures
- the use of non-reentrant library functions
- access to common variables, etc.


## Multithread programming models

- ❖ Three multithread programming models exist
  - User-level thread
    - Thread implemented at user-level
    - The kernel is not aware that threads exist
  - Kernel-level thread
    - Thread implemented at kernel-level
    - The kernel directly supports the thread concept
  - Mixed or hybrid solution
    - The operating system provides both user-level and kernel threads

## User-level threads

- ❖ The thread package is fully implemented in the user space, as a set of functions
  - These function calls, at user level, the standard system libraries
  - Each process has a its own thread table, which is managed by the thread package functions
  - The kernel is **not aware about** threads, it manages only processes

# User-level threads





## User-level threads

### ❖ Advantages

- Efficient management because it is carried out at the user level
- Fast context switching between threads of the same task
- Can be implemented on top of any kernel
- Allow the programmer to generate a large number of threads
- Possibly, you can implement your own custom scheduling strategy

## User-level threads

### ❖ Disadvantages

- There is only a single thread running per process task even in a multiprocessor system
- The scheduler maps all the user-level threads of a process to a the single kernel thread of that process
- If a thread makes a blocking system call, all threads of the same process are blocked
- There is no true parallelism


## User-level threads

### ❖ Disadvantages

- There is not automatic scheduling of the threads of a process
- Either a thread explicitly pass CPU control to other threads of the same process,
- or a scheduler is implemented, at user-level, using the kernel timer

## Kernel threads

- ❖ Threads are directly managed by the kernel



## Kernel threads

### ❖ Advantages

- A ready thread can be scheduled even if another thread of the same process has called a blocking system call
  - If a thread is blocked, the kernel scheduler can select another thread among the ready threads of all processes
- In a multiprocessor system multiple threads of a process can be executed in real concurrency

## Kernel threads


### ❖ Disadvantages

- Context switching more expensive because it requires passing through the kernel mode, and kernel support
- The maximum number of threads per task/system is limited, due to the use of kernel memory structures

## Hybrid implementation

- ❖ One of the multi-thread programming problems is to define the relationship between user threads and kernel threads
  - Virtually all modern operating systems have a kernel thread
    - Windows, UNIX, Linux, MAC OS X, Solaris
  - The basic idea is to have  $m$  user threads and to map them to  $n$  kernel threads
  - Typically,  $n < m$

## Hybrid implementation





## Hybrid implementation

- ❖ The hybrid implementation attempts to combine the advantages of both approaches
  - The user decides the number of its user-level threads, and the number of kernel-threads on which they must be mapped
  - The kernel is aware only of the kernel thread and only manages those threads
  - Each kernel thread can be used in turn by several user threads

## Processes and threads coexistence

### ❖ Several problems arise due to the coexistence of processes and threads

#### ➤ Using the system call `fork`

- A `fork` duplicates only the thread that makes the call, or all the threads of the process?

#### ➤ Example

- P has two threads T1 and T2
- T1 is waiting on a read, while T2 performs a `fork`
- Now we have P and his child, both with T1 waiting on a read
- Which thread T1 will receive data? Only one? Which? Both?

## Processes and threads coexistence

### ❖ Using the system call **exec**

- Does the **exec** replace only the calling thread with the new process, or all threads?
- Some systems provide different versions of the system call **fork**
- **forkall**, duplicates all process threads
- **fork1**, duplicates only the calling thread

### ❖ Signal management

- If a process receives a signal which thread will catch it?
- What happens if multiple threads indicate their interest to catch a signal? Which will handle the signal?