

The background of the slide features a large, stylized white 'C' logo on the right side, set against a blue background with abstract, swirling white lines. On the left side, there is a block of C code.

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>

#define MAXPAROLA 30
#define MAXRIGA 80

int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
                           delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE * f;

    for(i=0; i<MAXPAROLA; i++)
        freq[i]=0;

    if(argc != 2)
    {
        fprintf(stderr, "ERRORE: serve un parametro con il nome del file\n");
        exit(1);
    }
    f = fopen(argv[1], "rt");
    if(f==NULL)
    {
        fprintf(stderr, "ERRORE: impossibile aprire il file %s\n", argv[1]);
        exit(1);
    }

    while( fgets( riga, MAXRIGA, f ) != NULL )
```

Critical Sections – Mutual exclusion

Software solutions

Stefano Quer and Pietro Laface

Dipartimento di Automatica e Informatica

Politecnico di Torino

Requirements

- ❖ Any solution to the problem of the CS must meet the following requirements
 - Mutual Exclusion
 - One thread at a time can gain access to a Critical Section
 - Progress (no deadlock)
 - The solution allows a thread to access its CS, in a finite amount of time, even if another process is blocked **outside** its critical section.

Requirements

- Limited waiting (No starvation)
 - There must be a limited number of times that other threads are allowed to access their CS before a thread that made its access reservation is able to gain control of its CS.
- Any solution should be symmetric
 - The selection of which thread may access its CS should not depend on the
 - relative priority of the threads
 - relative speed of the threads

Software solution: no special instructions

- ❖ The software solutions to the CS problem are based on the use of shared variables
- ❖ We will analyze the solution with only two threads T_i and T_j
- ❖ The proposed solution is not easily extended to more than two threads

Mutual exclusion: Solution 1

❖ Shared variables

➤ `int flag[2] = {FALSE, FALSE};`

P_i / T_i

```
while (TRUE) {  
    while (flag[j]);  
    flag[i] = TRUE;  
    CS  
    flag[i] = FALSE;  
    non critical section  
}
```

P_j / T_j

```
while (TRUE) {  
    while (flag[i]);  
    flag[j] = TRUE;  
    CS  
    flag[j] = FALSE;  
    non critical section  
}
```

Mutual exclusion
Deadlock
Starvation
Symmetry

?

Mutual exclusion: Solution 1

❖ Shared variables

➤ `int flag[2] = {FALSE, FALSE};`

P_i / T_i

```
while (TRUE) {  
    while (flag[j]);  
    flag[i] = TRUE;  
    CS  
    flag[i] = FALSE;  
    non critical section  
}
```

P_j / T_j

```
while (TRUE) {  
    while (flag[i]);  
    flag[j] = TRUE;  
    CS  
    flag[j] = FALSE;  
    non critical section  
}
```

❖ Mutual exclusion **not** granted

➤ T_i and T_j can access to their CS at the same time

Mutual exclusion: Solution 1

❖ Solution

- A shared vector of flags "busy CS "
- A thread tests the other thread "busy CS " flag and sets its own

❖ It **does not guarantee** mutual exclusion in CS

❖ The technique fails because

- The lock variable is controlled and changed by two statements
- A context switching may occur between the two statements (they **are not** executed as single, **atomic** instruction)

Mutual exclusion: Solution 1

- ❖ Flag "Busy CS " is a **lock** variable
 - It serves to protect the CS
- ❖ Even if the solution were correct, the cycles testing the flag is a **busy form of waiting**
 - Waste of CPU time
 - Acceptable only if the busy wait is very short
- ❖ This lock mechanism, which uses the busy form of waiting, is called **spin lock**

Mutual exclusion: Solution 2

❖ Shared variables

➤ `int flag[2] = {FALSE, FALSE};`

Exchanges test and set statements

P_i / T_i

```
while (TRUE) {  
    flag[i] = TRUE;  
    while (flag[j]);  
    CS  
    flag[i] = FALSE;  
    non critical section  
}
```

P_j / T_j

```
while (TRUE) {  
    flag[j] = TRUE;  
    while (flag[i]);  
    CS  
    flag[j] = FALSE;  
    non critical section  
}
```

Mutual exclusion
Deadlock
Starvation
Symmetry

?

Mutual exclusion: Solution 2

❖ Shared variables

➤ `int flag[2] = {FALSE, FALSE};`

P_i / T_i

```
while (TRUE) {  
    flag[i] = TRUE;  
    while (flag[j]);  
    CS  
    flag[i] = FALSE;  
    non critical section  
}
```

P_j / T_j

```
while (TRUE) {  
    flag[j] = TRUE;  
    while (flag[i]);  
    CS  
    flag[j] = FALSE;  
    non critical section  
}
```

❖ Possible deadlock

➤ Both threads can set their flag to TRUE, and wait forever

Mutual exclusion: Solution 2

- ❖ Solution 2 tries to solve the problem with a dual approach
 - Reserves the access to the CS before testing its availability (i.e., performs setting before testing)
 - But deadlock is possible
 - Again, busy form of waiting with spin lock

Mutual exclusion: Solution 3

❖ Shared variables

➤ `int turn = i;`

Or
`int turn = j;`

P_i / T_i

```
while (TRUE) {  
    while (turn!=i);  
    CS  
    turn = j;  
    non critical section  
}
```

P_j / T_j

```
while (TRUE) {  
    while (turn!=j);  
    CS  
    turn = i;  
    non critical section  
}
```

Mutual exclusion
Deadlock
Starvation
Symmetry

?

Mutual exclusion: Solution 3

❖ Shared variables

➤ `int turn = i;`

Or
`int turn = j;`

P_i / T_i

```
while (TRUE) {  
    while (turn!=i);  
    CS  
    turn = j;  
    non critical section  
}
```

P_j / T_j

```
while (TRUE) {  
    while (turn!=j);  
    CS  
    turn = i;  
    non critical section  
}
```

❖ Does not comply with the requirements

- T_i and T_j access their CS only alternatively
- If T_i (T_j) has not interest in using its CS, P_j (P_i) cannot enter its CS (**starvation**)

Mutual exclusion: Solution 3

❖ Solution 3 uses

- A binary variable "turn", which indicates that the thread is enabled to enter its CS
- Mutual Exclusion is ensured by assignment of the access turn
- The solution involves alternation and possible starvation
- Busy form of waiting with spin lock

Mutual exclusion: Solution 4

❖ Shared variables

- `int turn = i;`
- `int flag[2] = {FALSE, FALSE};`

Or
`int turn = j;`

```
while (TRUE) {  $P_i / T_i$   
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] &&  
        turn==j);  
    CS  
    flag[i] = FALSE;  
    non critical section  
}
```

```
while (TRUE) {  $P_j / T_j$   
    flag[j] = TRUE;  
    turn = i;  
    while (flag[i] &&  
        turn==i);  
    CS  
    flag[j] = FALSE;  
    non critical section  
}
```

Mutual exclusion
Deadlock
Starvation
Symmetry

?

Mutual exclusion: Solution 4

❖ Shared variables

- `int turn = i;`
- `int flag[2] = {FALSE, FALSE};`

Or
`int turn = j;`

Mutual
exclusion?

```
while (TRUE) {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] &&  
        turn==j);  
    CS  
    flag[i] = FALSE;  
    non critical section  
}
```

P_i / T_i

In CS iff
`flag[j]==FALSE OR turn==i`

T_i and T_j both in their CSs?
No, because `turn==i` or `turn==j`,
not both

If T_j is in its CS, T_i can enter its CS?
If T_j is inside its CS, `flag[j]==TRUE` (set by T_j)
AND `turn==j` (set by T_j),
thus T_i will wait

Mutual exclusion: Solution 4

❖ Shared variables

- `int turn = i;`
- `int flag[2] = {FALSE, FALSE};`

Or
`int turn = j;`

Deadlock?

```
while (TRUE) {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] &&  
           turn==j);  
    CS  
    flag[i] = FALSE;  
    non critical section  
}
```

P_i / T_i

T_i/T_j wait only on this while loop

If T_i is waiting and T_j is not interested in its CS, `flag[j]==FALSE`, thus T_i can access its CS

T_i and T_j cannot be both waiting, because variable **turn** stores a **single value at a time**

If T_i is waiting and T_j releases its CS, T_j sets `flag[j]=FALSE`, thus T_i can access its CS

Mutual exclusion: Solution 4

❖ Shared variables

- `int turn = i;`
- `int flag[2] = {FALSE, FALSE};`

Or
`int turn = j;`

Starvation?

```
while (TRUE) {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] &&  
        turn==j);  
    CS  
    flag[i] = FALSE;  
    non critical section  
}
```

P_i / T_i

T_j is in its CS, and is very fast at reserving again access to its CS. Can T_i wait forever (starve)?

Testing (also) variable **turn** breaks the deadlock, which is possible in Solution 2

T_j sets `flag[j]` to FALSE but immediately after to TRUE. However, it sets `turn=i`, enabling access for T_i thus T_j will wait

Mutual exclusion: Solution 4

❖ Shared variables

- `int turn = i;`
- `int flag[2] = {FALSE, FALSE};`

Or
`int turn = j;`

Symmetric?

```
while (TRUE) {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] &&  
        turn==j);  
    CS  
    flag[i] = FALSE;  
    non critical section  
}
```

P_i / T_i

```
while (TRUE) {  
    flag[j] = TRUE;  
    turn = i;  
    while (flag[i] &&  
        turn==i);  
    CS  
    flag[j] = FALSE;  
    non critical section  
}
```

P_j / T_j

Symmetrically identical codes

Mutual exclusion: Solution 4

- ❖ The first software solution that allows two or more processes to share a single-use resource without conflict, using only shared memory and normal instructions, has been proposed by G. L. Peterson [1981]
 - Solution with overhead, and busy form of waiting

Conclusions

- ❖ In general, the software solutions to the problem of CS are complex and inefficient
 - Setting and testing a variable by a thread is an operation that is "invisible" to the other threads
 - **Test and set operations are not atomic**, thus a thread can react to the presumed value of a variable rather than to its current value
 - The solutions for n threads are even more complex